

18-03

40:02 – 40:06

so it's a it should now be installed into the system

So, 它现在应该被落地到系统中了

40:11 – 40:12

all right

40:14 – 40:14

So okay

40.14–40.15

so for the rest of this lecture

So, 在这节课剩下的时间里

40.15–40.18

,let's just assume that we're always doing forward validation

假设, 我们使用的始终是forward validation

40:18 – 40:19

so anytime you want to validate a transaction

So, 每当你想对一个事务进行验证时

40.19–40.21

you find all newer transactions than you

你会找到所有比该事务时间戳更大的事务

40.21–40.23

and you perform this validation step

你就会执行这个验证步骤

OCC – VALIDATION STEP #1

T_i completes all three phases before T_j begins.

40:26 – 40:29

okay,so there's a couple of different scenarios that we have to cover

So, 我们需要介绍两种不同的情况

40:29 – 40:30

okay the first scenario is that

Ok, 第一种情况是

40.30–40.32

if you find a transaction T_j


如果你找到一个事务 T_j

40:34 – 40:42

~~where all your phases your read your would sorry, your work your that your validation in~~

your write steps happen before anything that T_j has done

你原有事务中的写操作要在 T_j 这个时间戳更大的事务执行任何操作前完成



40:42 – 40:45

all right this is sort of like the simple happy path case here here

这就如这里这个简单的例子所展示的

40:45 – 40:47

so this is an example

So, 在这个例子中

40:47–40:48

you have t1 and t2

我们有两个事务，即T1和T2

40:48–40:52

,t1 completes all of its steps before t2 is then anything

在T2执行任何操作之前，T1就完成了它的所有步骤

40:53 – 40:55

okay so this is a really simple example

Ok, 这是一个非常简单的例子

40:55–40:58

you could essentially collapse this into one and now you have a serial execution

本质上来讲，你可以将这些事务合并为一个，然后你就得到了一个Serial execution

40:59 – 41:00

so you dont only have to do anything

So, 你无需去做任何事情

41:00–41:04

all that interesting here this is sort of given to you naturally

这里所有你感兴趣的東西都是默认很自然的提供给你的

OCC – VALIDATION STEP #2

T_i completes before T_j starts its **Write** phase, and
 T_i does not write to any object read by T_j .
 $\rightarrow \text{WriteSet}(T_i) \cap \text{ReadSet}(T_j) = \emptyset$

41:06 – 41:11

the second scenario that we have to handle, is that

我们需要处理的第二种情况是

41:11–41.17

if ~~t1 completes before Tj~~, sorry T1 completes before Tj begins its write phase

如果Tj在开始执行它的写阶段之前，Ti就完成了它的工作

41:18 – 41:19

okay

41.19–41.21

and we have to make sure in this scenario

我们需要确保，在这种情况下

41.21–41.28

that the stuff that we write the the write set in our transaction ,doesn't intersect with the read set of the other transaction

我们事务中写操作要处理的东西不会与其他事务读操作涉及的东西相交

41:28 – 41:32


so the other transaction hasn't read anything that we're going to that we've written

即其他事务不会去读取任何我们要写入的东西

41:32 – 41:33

okay why is this important

为什么这点很重要呢？



41.33–41.35

so let's walk through an example

So, 我们来看个例子

41:35 – 41:38

so in this scenario we have two transactions

So, 在这个场景下, 我们有两个事务

41.38–41.39

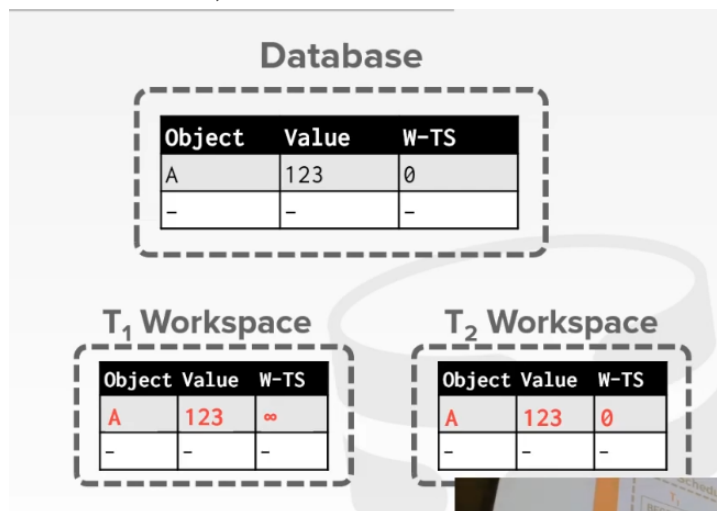
t_1 is reading and writing a

T_1 对A进行读取和写入操作

41.39–41.42

t_2 is reading a and doing nothing else

T_2 除了读取A以外, 其他什么事也不做



41:42 – 41:46

so so this is the state of the database

So, 这就是数据库的状态

41.46–41.50

so t_1 has read a and written to it assume that is writing the same value 123

So, T_1 对A进行读取, 并将它的值修改为123

41:51 – 41:52

and it's timestamp is infinity,

它的时间戳是无穷大

41.52–41.54

T2 is read a

T2对A进行读取

41.54–41.56

so it has a timestamp of 0 there

So, 它的时间戳就是0

41.56–41.59

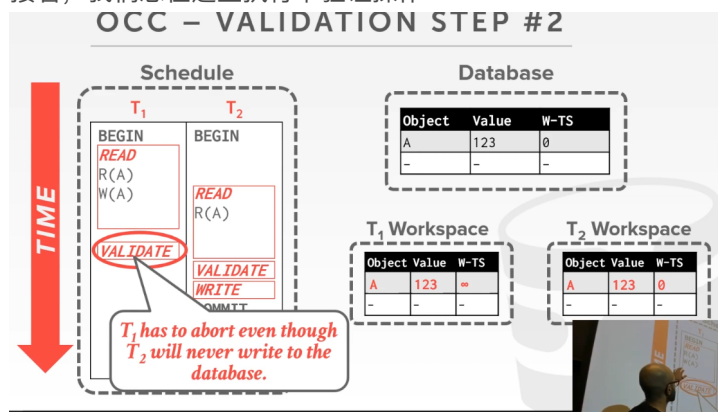
,and it has a local copy in its private workspace

在它的私有工作空间中, 它保存了A的一份本地副本

41:59 – 42:02

and then you want to perform the validation here

接着, 我们想在这里执行下验证操作



42:03 – 42:05

right so t1 has to abort

So, 这里需要中止T1

42:05–42:06

because it's read set

图标为它是read set

42:06 – 42:09

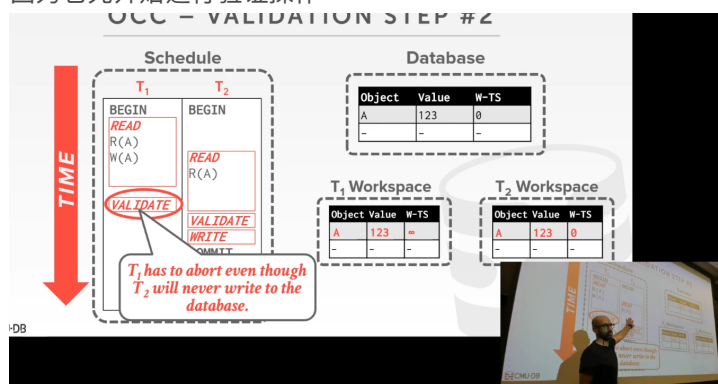
so it has a lower timestamp than t2

So, 它的时间戳要比T2小

42:09–42.11

,because it begins validation first

因为它先开始进行验证操作



42:11 – 42:14

So it's write set intersects with the read set here

So, 它的write set与此处的read set相交了

42:15 – 42:17

all right so that violates the invariant that we had before

So, 这违反了我们之前拥有的不变量

42.17–42.21

so we have to abort this transaction

So, 我们需要中止这个事务 (T1)

42:22 – 42:23

okay

42.23–42.24

in the serial order

在Serial Order的情况下

42.24–42.26

t1 has to begin before t2,

T1必须在T2之前开始

42.26–42.28

but t2 has read a stale value

但T2读取到了一个过时的值

42:29 – 42:32


so we have to abort

So, 我们需要中止事务

42:34 – 42:37

I think it's a pretty simple example

我觉得这是一个相当简单的例子



42:39 – 42:42

But now let's that slightly tweak it

现在让我们来稍微调整下这个例子

42.42–42.46

and now we'll have T2 begin validation before t1

我们让T2在T1之前先执行验证

42:46 – 42:50

so in the serial order t2 happens before t1

So, 在Serial Order中, T2在T1之前执行

42:51 – 42:53

what's gonna happen in this in this case right

在这个例子中, 会发生什么呢?

42.53–42.58

,so t1 reads a ,begins validation, it doesn't have a write set

So, T1对A进行读取, 接着开始验证, 这里它并没有write set

42:58 – 43:01

so it doesn't it doesn't have anything to intersect with this transaction

So, 它也就没有任何与这个事务相交的部分

43.01–43.02

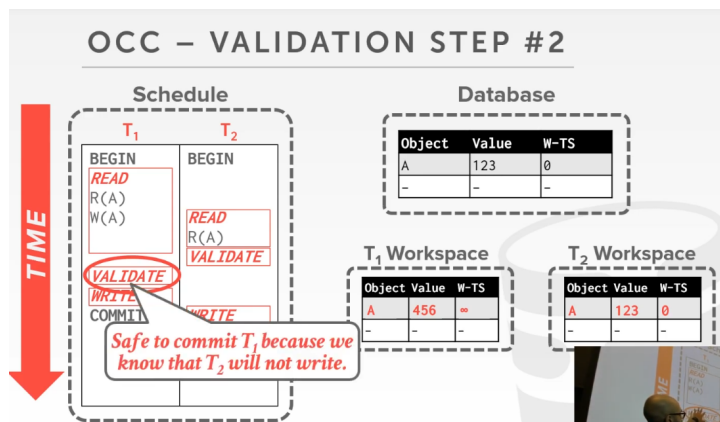
,so so it commits

So, 我们可以提交T1

43:03 – 43:04

t2 does the same thing

T2也做了相同的事情



43:09 – 43:17

it's write set ,doesn't intersect with the the read set of T, it doesn't intersect with any other concurrent transaction

T2的write set不与其他并发事务中的操作相交

43:18 – 43:20

because it's what this one's all revalidated

因为它已经被重新验证了

43:21 – 43:23

so ~~another thing at like~~ nothing happens right

So, 这里什么也没发生

43:23 – 43:23

So when the serial order

So, 当我们是Serial Order的情况下

43.23–43.24

t2 happens before t1

T2在T1之前开始执行

43.24–43.28

it reads the value of a that initially existed before

它读取了A之前一开始就存在的值

43:28 – 43:31

This transaction is going to read the value of a there

该事务就会去读取A的值

43.31–43.34

make a local copy read write to it,

它会去制作A的一份本地副本, 并对该副本进行读写操作

43.34–43.34

and everything's okay

所有东西都是ok的

43:39 – 43:43

all right,questions, yes

有问题吗? 请问

43:50 – 43:51

so in this scenario

So, 在这种情况下

43.51–43.54

because T2 begins validation first

因为T2先开始进行验证

43.54–43.56

it appears first in the serial order

它会先出现在Serial Order中

~~44:04 – 44:10~~

~~yeah yeah~~

44:16 – 44:20

so T1 has to abort

So, 我们需要中止T1

44.20–44.24

because in the serial order what should have happened is that

因为在Serial Order中, 这里应当会发生的事情是

44.24–44.27

T2 because it begins validation first

因为T2先开始进行验证

44.27–44.29

, it's the first to recognize that there's something wrong

它会首先注意到这里出了问题

44:30 – 44:36

so it recognizes something's wrong as soon as it begins a validation

So, 只要当它开始验证时, 它就会意识到这里有东西出了错

44:36 – 44:37

because it gets assigned a timestamp,

因为它被分配了一个时间戳

44.37–44.40

it sees a concurrent transactions that's read stale value

它看到了一个并发的事务, 该事务读取到了一个过时的值

44.40–44.41

so it aborts

So, 该事务被中止了

44:45 – 44:47

well when T2 goes to validate,

Well, 当T2开始验证的时候

44.47–44.49

T1 note is no longer a valid transaction

T1被打上标记, 它不再是一个有效的事务

44:49 – 44:51

So it won't see any concurrent transactions,

So, 它不会看到任何并发的事务

44.51–44.55

there's nothing to intersect with also it's a read-only transaction

这样也就没有其他事务会与其相交, 即便它是一个只读事务

44:55 – 44:56

so there's nothing to intersect with

So, 这里不会有任何事务与它相交

45:16 – 45:22

~~T2 couldn't~~, no not until the stalls it's it's changes

No, T2不会去暂停它所做的修改

45:20-45:22

because it's a private private thing

因为它是一个私有的东西

45:25 - 45:28 - 45:28

yeah yeah yeah question

请讲

45:54 - 46:00

~~Yeah so when~~ when this T1 begins its validate phase

你是说当T1开始执行它的验证阶段的时候?

46:00 - 46:02

you don't get a timestamp

你不会拿到时间戳

46:02-46:03

yeah I don't get I don't know 1 or 2 or something

你不会拿到1或者2之类的时间戳

46:04 - 46:07

but this T2 doesn't have a write timestamp yet

但T2目前还没有write timestamp

46:08 - 46:09

yeah I mean it's infinity.

我的意思是, 它是无穷大的

46:09-46:11

we doesn't have any timestamp

我们还没有分配任何时间戳

46:11-46:12

because it hasn't begun validation yet

因为它还未开始验证操作

46:15 - 46:18

no no no so no no no no

No No No, 你讲的不对

46:19 - 46:21

so the timestamp is at assigned to the transaction

So, 我们会将时间戳分配给这个事务

46:21-46:30

~~it isn't assumed that~~ it doesn't take whatever timestamp is from the the tuple

我们并不会去使用tuple所携带的任何时间戳

46:30 - 46:32

yeah so that's just a local copy

So, 这只是一个本地副本

46:33 - 46:34 - 46:40 - 46:41

okay yeah yeah yeah

请说

46:56 - 46:57

are you saying

你说的是不是这么一回事

46:57-47:01

if t2 performs a write to this object A

如果T2对A执行写操作

47:02 – 47:08

it'll modify the write timestamp of its local copy to infinity

它会将它本地副本上的write timestamp修改为无穷大

47:10 – 47:10

yeah

47:25 – 47:26

Yeah so if it makes a write

So, 如果它执行了这个写操作

47:26–47:29

~~it will get a local time~~ it'll get a timestamp in' validation

它会在验证的时候拿到一个时间戳

47:29–47:34

and it'll update the write timestamp there, If there's no conflicts yeah

如果这里没有任何冲突的话, 它会更新这里的write timestamp

47:58 – 47:58

no no

No No

47:58–48:00

you have to have this write be atomic

你所拥有的这个写操作必须是原子性的

48:00–48:04

so you write all all your changes back into into the database in this phase

So, 你会在这个阶段将你所做的所有修改都写回数据库

48:11 – 48:14

that's just because of the schedule that we have here

这只是因为我们这里所拥有的这个schedule

48:14 – 48:17

~~so we have the device~~ you can think of it as like we begin the validation

So, 你可以想象一下, 当我们开始验证的时候

48:17–48:19

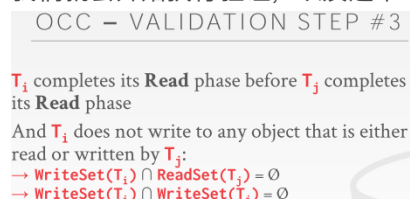
we finish and do context switch here

当我们执行完毕, 进行上下文切换到T1这里的时候

48:19–48:28

and now we begin the validation and the write phase

我们就会开始执行验证, 以及这个写阶段



48:28 – 48:28

all right

48:28–48:30

so this is the second example

So, 这是第二个例子

48.30–48.33

there's still one more to go

我们还要再看一个例子

48:34 – 48:34

in this one

在这个例子中

48.34–48.44

the transaction you have to ensure that the transaction T₁ completes its read phase

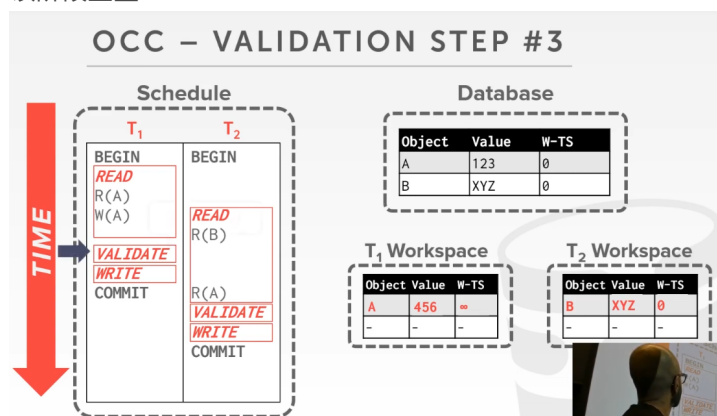
before transaction T₂ begins its read phase

你需要确保在T₂开始它的读阶段开始之前，T₁完成了它的读阶段

48:44 – 48:55

so we have to ensure that the write set of my transaction doesn't intersect with the the read set and the write set of all these of the transactions that fall into this category or the read phases are overlapping

So, 我们需要确保我们事务中的write set不与其他事务的read set和write set相交, 或者不与读阶段重叠



48:56 – 49:01

and this is another example of how we can go through this

这是我们要看的另一个例子

49:02 – 49:04

so let's see in this scenario

So, 在这个场景下

49.04–49.06

t₁ is reading and writing a

T₁对A进行读写操作

49:07 – 49:08

oops

不好意思, 按错了

49.08–49.10

t₂ is reading B and reading a,

T₂先读取B, 再读取A

49.10–49.14

t₁ and begin that validation phase

T₁开始执行它的验证阶段

49:14 – 49:16

so it gets a timestamp of 1

So, 它拿到的时间戳是1

49.16–49.19

,and now I can commit it

我可以提交T1

49.19–49.27

,because my write set doesn't intersect with the read set of transaction t2 at this point in time

因为T1中的write set此时并没有与T2中的read set相交

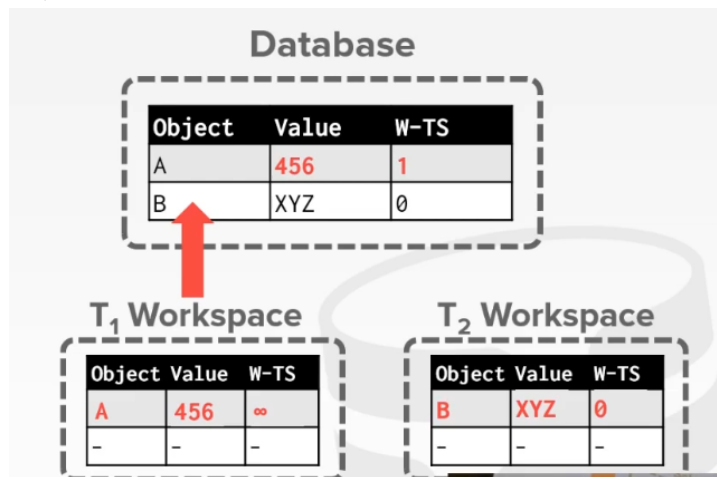
49:29 – 49:30

okay

49.30–49.34

so I can install my changes back into the into the database

So, 我可以将我所做的修改落地到数据库中



49:35 – 49:38

I update 4 5 6 or update A with value 4 5 6

我将A的值更新为456

49.38–49.39

and I set my write timestamp to 1

并将它的write timestamp设置为1

49.39–49.45

which is the timestamp that I was assigned at validation phase

这也是我在验证阶段所分配的时间戳

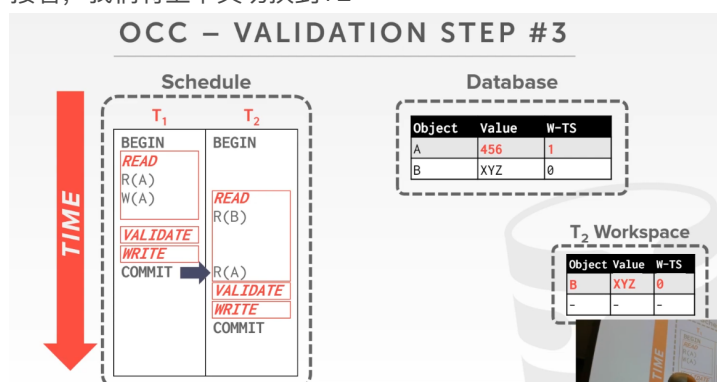
49:45 – 49:45

okay

49.45–49.51

then when we come back down into we do a context switch back into t2

接着, 我们将上下文切换到T2



49:51 – 49:54

t2 does a R(B) or sorry a R(A)

接着，T2对A进行读取

49.54–49.58

,and now it can go back to the main it goes to the to the main database

现在，它跑到主数据库那里

49:58 – 50:03

and does a local copy of the updated value of A made by t1

读取到由T1更新过的A的值的本地副本

50:04 – 50:06

and then it can validate and write successfully

接着，它可以进行验证，并成功写入数据

50.06–50.08

,because there are no concurrent transactions

因为这里并没有有任何其他的并发事务

50:11 – 50:13

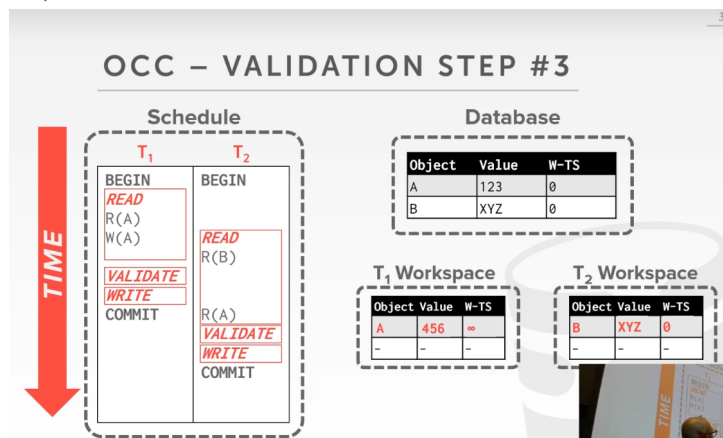
Good yeah

都懂了吧，请讲

50:22 – 50:26

so if you let's go back to this example

So，我们回到这个例子



50:28 – 50:30

so when you want to perform a write

So，当你想执行一次写操作的时候

50.30–50.36

right you make a local copy of the database object into your local private workspace

你会制作该数据库对象的一份副本，并保存到你的本地私有工作空间中

50:36 – 50:38

and then that's why you set the timestamp to infinity

这就是你将时间戳设置为无穷大的原因所在

50:38 – 50:39

you set it to infinity

你将它设置为无穷大的原因是

50.39–50.41

,because you don't know what your timestamp is going to be

因为你不知道你的时间戳会是什么

50:41 – 50:44

because you haven't been assigned one yet, right

因为你还未分配这个时间戳

50:45 – 50:49

it's zero here only for mainly for illustrative purposes

这里我们之所以写0，只是出于展示原因

50:49–50:55

, you can think about like there was some transaction with timestamps 0 ,that bulk loaded all of these objects into the database

你可以想象一下这种场景，某个时间戳为0的事务将所有的对象加载到数据库中

50:55 – 50:57

so that's why the write timestamp is 0

So，这就是为什么write timestamp是0的原因了

51:00 – 51:00

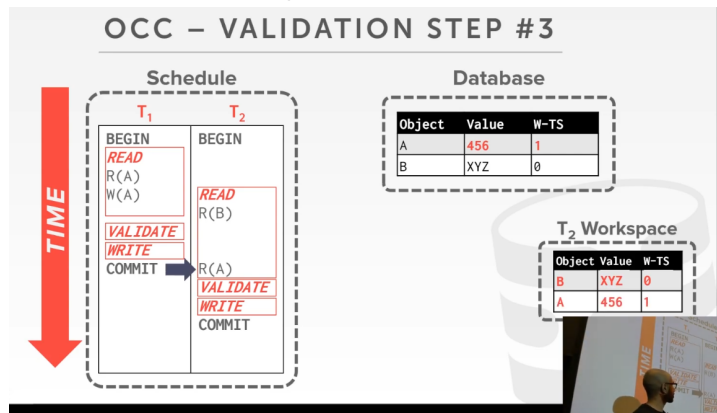
when you read it

当你读取它的时候

51.00–51.11

you read it whatever timestamp is associated with the object at that point

该对象的read timestamp就与你读取它的时间相关



51:11 – 51:14

yeah so t2 does a R(A)

So, T2执行了R(A)

51.14–51.16

does a valid does a validation does a write

并执行了验证以及写操作

51.16–51.17

everything is good

这一切都没问题

51.17–51.18

because there's no concurrent transactions

因为这里并没有其他的并发事务

51:19 – 51:19

yeah

请讲

51:23 – 51:26

yeah these are all for validation yeah correct

这些都是为了进行验证，说的没错

51:34 – 51:37

yeah you still need to rely on the same timestamps

没错，你依然需要依靠这些相同的时间戳

OCC – OBSERVATIONS

OCC works well when the # of conflicts is low:

- All txns are read-only (ideal).
- Txns access disjoint subsets of data.

If the database is large and the workload is not skewed, then there is a low probability of conflict, so again locking is wasteful.

51:44 – 51:53

yeah so that was essentially the I'm cup we've been kind of hind wavy on a couple of things on, how you actually install the write atomically, and how you do the validations in a parallel way

So, 我们已经讲了两个东西, 即我们该如何保证我们的写操作原子性, 以及我们该如何以并行的方式进行验证

51:53 – 51:54

because in a real system

因为在一个真正的系统中

51:54–51:56

you have to do this parallel validation

你需要做这种并行验证

51:56 – 51:59

but I think hopefully the main idea has gotten across right

但我希望你们已经理解了这里面的主要思想

52:00 – 52:01

it's not too complicated

它并不是很复杂

52:01–52:03

there's some trickeries here

这里面有些棘手的地方

52:03–52:06

but generally it's it's pretty intuitive

但总的来讲, 它还是比较直接的

52:07 – 52:10

so there's a few observations that we should try to make

So, 我们从中学到了一些东西

52:11 – 52:18

all of these timestamp in optimistic concurrency protocols sort of work well when there's very little conflict right

当没有什么冲突的情况下, 乐观并发协议中这些基于时间戳的协议的效果非常好

52:18 – 52:25

because you allowed the transactions to proceed without acquiring locks without doing that anything heavyweight it worked

因为这允许我们在不获取lock以及不做那些代价很高的操作的情况下, 我们能够去处理这些事务

52:25–52:30

and at the end you do a sort of lightweight semi lightweight validation to make sure that your transaction is still valid

在最后, 我们会做些轻量级或者说半轻量级的验证来确保我们的事务是有效的

52:31 – 52:35

so if you have very few very few conflicts

So, 如果你执行的事务中几乎没有什么冲突

52:35–52:41

even better if all your transactions are essentially just doing read-only work

更好的情况是，如果你的事务所做的都是只读型工作

52:41 – 52:43

and if the access disjoint sets of data

并且它们访问的都是不相交的数据集

52.43–52.45

then this these protocols work very well

那么，这些协议的效果就会非常好

52.45–52.48

because they're actually not doing much work at all aside from doing some local data copying

因为它们除了会去制作本地数据副本以外，它们并不会做太多工作

52:49 – 52:53

Right, another way to think about this is that

思考它的另一种方法是

If the database is large and the workload is not skewed, then there is a low probability of conflict, so again locking is wasteful.

52.53–52.57

if you have a very large workspace and very small transactions

如果你的数据库很大，并且你要执行的事务数量也不多

52.57–53.02

and the probability of these transactions overlapping in the read and write set is very low

同时，这些事务在read set和write set上重叠的可能性又比较低

53:02 – 53:03

in this specific scenario

在这种特定情况下

53.03–53.06

these optimistic concurrency protocols work very well

这些乐观并发控制协议的效果会非常好

53:07 – 53:08

on the other hand

另一方面

53.08–53.11

in very highly contentious workloads

在那些存在着很多冲突的workload中

53:11 – 53:12

what ends up happening is that

最终会发生的事情是

53.12–53.14

you have transactions consistently restarting right

你的事务始终会反复重启

53:14 – 53:19

so they do a lot of work in their work phase

So，在它们执行阶段，它们会做大量的工作

53:20 – 53:22

assuming if there is no contention
假设，如果这里面没有冲突的情况发生

53.22–53.27

and then and only at the very end do they figure out oh crap all this work that I've done
is kind of useless after we start

只有到了最后，它们发现：Oh，自我开始执行后，我所做的所有工作都是无谓的消耗

53:27 – 53:33

right so you're kind of deferring a lot of the heavy lifting towards the end

So，你就会将很多繁重的工作都放到最后去做

53.33–53.35

assuming that there won't be a lot of heavy lifting

假设这里并没有那么多繁重的工作

53:35 – 53:38

but in the contentious work though there is heavy lifting

但在冲突很多的workload中，这里面就存在着那些繁重的工作

53.38–53.41

because your sense they have to abort everything you've done

因为它们必须中止你已经完成的所有东西

53:41 – 53:44

so really I think the the the research sort of shows that

So，这里的研究表明

53.44–53.46

in contentious workloads

在这些冲突很多的workload中

53.46–53.53

both two-phase locking and optimistic concurrency protocols generally don't work,

两阶段锁和乐观并发协议通常来讲都不可行

53.53–53.54

they're sort of like almost equivalent

它们的效果几乎是一样的

53:54 – 53:58

so when **2PL** ~~simply acquiring there's~~ there's a lot of contention on locks

So，在两阶段锁中，它里面存在着大量争抢锁的情况

53:59 – 54:02

so you have transactions waiting and these hot locks

So，你需要让事务去等待获取这些热门的lock

54.02–54.07

here you have transactions that are doing all their work almost wasteful work ,and then
aborting at the very end

这里你会遇上那些做了很多无用功的事务，它们在最后的时候被中止

54:07 – 54:10

so neither really work for contentious workloads

So，对于那些存在着冲突的workload来说，该协议的效果并不好

54:11 – 54:13

but if your workload is it has very low contention

但如果你的workload中所存在的冲突数量并不多

54.13–54.17

, then in general these optimistic protocols work better

那么，通常来讲，使用这些乐观协议的效果会很不错

54:17 – 54:20

because they have no overhead

因为它们并没有什么开销

54.20–54.23

and very very minimal override aside from some local data copy

并且它们对本地数据副本也只做最小程度的覆写

54:23 – 54:25

okay whereas in the 2PL protocol

然而在两阶段锁中

54.25–54.26

you actually have to acquire locks

实际上，你必须去获取这些lock

54.26–54.30

, even if there's no logical reason that you should be acquiring these locks

即便没有什么逻辑上的理由，你也应该去获取这些lock

54:36 – 54:36

Yeah

54:53 – 54:56

so the timestamps are assigned at validation

So，我们会在验证的时候分配这些时间戳

54:57 – 54:57

yeah

55:19 – 55:22

I think what you're describing is essentially basic timestamp order ,right

我觉得你所描述的本质上来讲就是Basic Timestamp Ordering

55:23 – 55:27

so in basic timestamp ordering, you assign the timestamps when the transaction begins

So，在Basic Timestamp ordering中，你会在事务开始的时候分配时间戳

55:27 – 55:28

In OCC

在OCC中

55.28–55.34

the protocol dictates that you assign the timestamps at validation time

该协议会让你在验证阶段的时候分配时间戳

55:34 – 55:39

because you want to defer the work of doing the actual checking to when you're ready to commit

因为你会将这些实际的检查工作延迟到你准备提交事务的时候再做

55:39 – 55:46

whereas basic timestamps ordering is saying ,every operation I'm going to perform, I

want to do a check to make sure that this is a valid operation

Basic timestamp ordering表示，我会对我要执行的每个操作进行检查，以确保这是一个有效操作

55:47 – 55:48

so just how the protocol works

So，这就是该协议的工作方式

55:50 – 55:50

yeah

55:58 – 56:04

you still need to make a local copy even if you're reading to ensure repeatable reads
即使你是在进行读操作的时候，你也依然需要制作一份本地副本以支持可重复读

OCC – OBSERVATIONS

OCC works well when the # of conflicts is low:

- All txns are read-only (ideal).
- Txns access disjoint subsets of data.

If the database is large and the workload is not skewed, then there is a low probability of conflict, so again locking is wasteful.

56:11 – 56:18

right so this slide is essentially just saying that for low contention workloads, these optimistic concurrency protocols work very well

So, 这张幻灯片上的内容是，当你workload中存在的冲突数量不多时，这些乐观并发协议的效果就会很好

56:18 – 56:19

because they have very little overhead

因为它们几乎没有什么开销

56:19–56:22

in comparison to two-phase locking

比起两阶段锁来说

56:22–56:26

which you know even if you have disjoint working sets between transactions

即使你在事务间的working set是不相交的

56:26 – 56:28

you still have to have transactions require locks

你依然会要求这些事务去获取锁

56:28–56:31

,whereas here there's no locks at all

但如果是OCC的话，就根本不需要用到锁

56:45 – 56:47

this guy yeah you're right

你说的没错

56:47 – 56:50

so you don't have to acquire like a database of a lock

So, 你不需要去获取数据库中的lock

56:50–56:51

, you still have to require latches

你依然需要去获取latch

56:51–56:56

, you still have to make sure the integrity of the data structure is valid

你依然需要去确保该数据结构的完整性是有效的

56:56 – 56:59

Right you can't just blindly overwrite stuff in memory

你无法盲目地去覆写内存中的数据

56.59–57.02

,because there could be concurrent access to these structures

因为这可能存在着并发访问这些结构的情况

OCC – PERFORMANCE ISSUES

High overhead for copying data locally.

Validation/Write phase bottlenecks.

Aborts are more wasteful than in 2PL because they only occur after a txn has already executed.

57:02 – 57:06

so I think you're kind of hitting on the point that we will talk about it I think in this slide

So, 我觉得你讲到我们要讲的重点了, 应该是在这个幻灯片上

57:08 – 57:12

yeah so there is a little bit of work overhead in OCC right

So, 使用OCC的时候, 我们会有一些开销

57:12 – 57:15

because you have to maintain local copies of everything that you want to read and write

因为你需要去维护你想要进行读写操作的那些对象的本地副本

57:15 – 57:16

so if you're updating a billion tuple

So, 如果你要更新十亿个tuple

57.16–57.17

you have to make a billion copies

你就必须制作十亿个副本

57.17–57.20

every transaction has to make a billion copies

每个事务都必须制作十亿个副本

57:20 – 57:22

and so there's a lot of overhead here

So, 这里面就存在着大量的开销

57:23 – 57:27

the other thing is that the validation and the write phase is, they were happening serially right

另一件事情是, 验证阶段和写阶段会按顺序执行

57:28 – 57:30

the only one transaction can be validating at a time

一次只会有一个事务进行验证

57:31 – 57:33

in reality you know in real systems

在真正的系统中

57.33–57.35

you have parallel validation and parallel writing

你可以进行并行验证和并行写

57:35 – 57:39

but again even those in those scenarios these phases become a big bottleneck

但在这些情况下, 这些阶段会变成一个巨大的瓶颈

57:41 – 57:42

and lastly

最后

57.42–57.45

because we're being optimistic

因为我们使用的是乐观方案

57.45-57.51

we're assuming that we can perform all of the work with ,you know safely without being interrupted by other transactions

我们假设我们执行的所有操作都不会被其他事务所打断

57:51 – 57:53

But in contentious workloads

但在那种存在冲突的workload中

57.53-57.54

that's not the case

这就不会是这种情况了

57.54-57.55

right that assumption is invalidated

这种假设就是无效的了

57:55 – 57:58

so in contentious workloads routing all this work upfront

So, 在这些存在着冲突的workload中, 我们会将所有的工作放在眼前

57:59 – 58:01

and then we could find out later that we actually have to abort

然后, 我们就能找出我们需要中止的那些工作

58:01 – 58:03


so all of this work is wasted

So, 所有这些被中止工作做的都是浪费的无用功

58.03-58.06

we're being optimistic and we're being wrong about that optimism

我们使用的是乐观方案, 但我们对这种乐观的看法是错误的



OBSERVATION

When a txn commits, all previous T/O schemes check to see whether there is a conflict with concurrent txns.
→ This requires latches.

If you have a lot of concurrent txns, then this is slow even if the conflict rate is low.

58:11 – 58:13

so when a transaction commits

So, 当一个事务要提交的时候

58.13-58.21

I think that the student was pointing this out correctly

我认为之前有学生正确地指出了这一点

58:21 – 58:25

even if the transactions logically don't overlap with each other

从逻辑上来讲, 即使这些事务彼此不会重叠

58:26 – 58:31

you still have to make sure that you maintain the physical integrity of the data structure

you're gonna be looking at

你依然必须确保你要去维护你所查看的数据结构的物理正确性(知秋注: 不会因为并发操作, 使程序中的一个对象内部数据状态发生异常)

58:31 – 58:34

so as part of the validation phase

So, 在执行验证阶段的时候

58.34–58.40

you have to look at oh I have you know all these other thousand transactions, I have to look inside read and write sets in a consistent way

你会看到所有这些数以千计的事务，我不得不深入其中通过一种可以保证一致性的操作方式来查看read和write set

58:41 – 58:45

I make this consistent by acquiring latches right

我通过获取latch来让它们保持一致

58:46 – 58:49

so this latch overhead it can actually play a big role

So，这种获取latch的开销实际上扮演了一个重要角色

58.49–58.52

even if they're logically just disjoint

即使从逻辑上来说，它们是不相交的

58.52–58.55

physically they're still contending on the same data structures

从物理上来讲，它们依然在抢夺同一个数据结构

58:55 – 58:59

right there so contending on the read and write data sets, even though logically they're disconnected

即使在逻辑上read set和write set是不相交的，但它们依然存在着争抢问题(知秋注:如果并发验证，那就会可能同时有多条线程查看同一个事务中本地保存的数据，这就会需要latch)

59:00 – 59:01

yes

请讲

59.01–59.03

so when I'm doing my validation right

So，当我在执行验证的时候

59:04 – 59:06

I have one transaction and me ,I'm ready to validate

我有一个准备进行验证的事务

59.06–59.09

I have to go into his read and write transaction set to make sure that we don't intersect

我必须去查看它的read set和write set，以确保它们是不相交的

59:10 – 59:10

all right

59.10–59.11

but he could be modifying it

但它可以对数据进行修改

59.11–59.14

because he's still running ,he or she's still running right

因为这个事务依然在执行

59:14 – 59:15

so I have to acquire latch

So，我需要去获取latch

59.15–59.20

read a consistent view of it , perform the intersection to make sure I'm okay

并读取到该数据的一个一致的状态

以这种一致性的方式来对它操作，以保证可以执行的很OK

59:20 – 59:24

and this other transaction can't modify that that said while I'm reading it

当我正在读取这个数据的时候，其他事务无法对该数据进行修改

59:29 – 59:34

but I'm reading the local copy, I'm reading that transactions local copy of the working set

你说的是，如果我读取的是该事务自己的本地副本数据？

59:37 – 59:44

Because the table doesn't have all of the updates that have been applied that this

transaction has read and written to write

因为数据库表中并没有更新该事务所做的所有更新

59:44 – 59:46

the transaction has a local copy of reads and writes

该事务对本地副本进行了读写操作

59:46 – 59:50

so to see if this transaction has read a value that I'm writing to

So，为了弄清楚该事务读取到的值是否是我正在写入的那个值

59:51 – 59:53

I only know that by looking at his read write set

我只有通过查看该事务的read set和write set，我才能知道这个

59:54 – 59:55

I don't know that by looking at the table

看数据库表，我是看不出这些的